

ECE194N HW 1

KNN Report

February 8, 2019

Student: Erik Rosten
Perm Number: 7143571
Email: erosten@ucsb.edu

Department of Electrical and Computer Engineering, UCSB

Code

My code for this programming problem is below

```
1 import numpy as np
2 import cifar_loader
3 import matplotlib.pyplot as plt
4 import os
5
6
7
8 def split_sets(all_images, all_classes, labels_to_keep):
9     images = []
10    classes = []
11    for i in range(labels_to_keep.shape[0]):
12        indices = np.where(all_classes == labels_to_keep[i])[0]
13        images.append(all_images[indices])
14        classes.append(all_classes[indices])
15
16    images = np.array(images)
17    images = images.reshape([-1,32,32,3])
18    classes = np.array(classes)
19    classes = classes.reshape([-1,])
20    return images, classes
21
22 def split_class_names(all_class_names, class_numbers):
23     class_names = []
24     for i in range(class_numbers.shape[0]):
25         class_names.append(all_class_names[class_numbers[i]])
26     return np.array(class_names)
27
28 def view_examples(train_imgs, train_cls, labels_to_keep, class_names):
29     print(class_names)
30     plt.figure(1)
31     for i in range(labels_to_keep.shape[0]):
32         label_class_name = class_names[i]
33         rand_integer = np.random.randint(0,4999)
34         # grab a random example index
35         index = np.where(train_cls == labels_to_keep[i])[0][rand_integer]
36         img_example = train_imgs[index]
37
38         ax = plt.subplot(1,labels_to_keep.shape[0],i + 1)
39         ax.axis('off')
40         imgplot = plt.imshow(img_example)
41         ax.set_title(label_class_name)
42     plt.show()
43
44
45 def rgb2gray_average(rgb_imgs):
46     gray_imgs = np.zeros((rgb_imgs.shape[0],32,32))
47     for i in range(rgb_imgs.shape[0]):
48         r, g, b = rgb_imgs[i,:,:0], rgb_imgs[i,:,:1], rgb_imgs[i,:,:2]
49         gray_imgs[i,:,:] = (r + g + b) / 3
50     return gray_imgs
```

```

51
52 def rgb2gray_luminosity(rgb_imgs):
53     gray_imgs = np.zeros((rgb_imgs.shape[0],32,32))
54     for i in range(rgb_imgs.shape[0]):
55         r, g, b = rgb_imgs[i,:,:0], rgb_imgs[i,:,:1], rgb_imgs[i,:,:2]
56         gray_imgs[i,:,:] = 0.21 * r + 0.72 * g + 0.07 * b
57     return gray_imgs
58
59 def euclidean_dist(img1, img2):
60     return np.linalg.norm(img1-img2)
61
62 def ssd(img1, img2):
63     return np.sum((img1-img2)**2)
64
65 def compute_dists(train_imgs, test_imgs):
66     num_test = test_imgs.shape[0]
67     num_train = train_imgs.shape[0]
68     dists = np.zeros((num_test, num_train))
69     for i in range(num_test):
70         print(i)
71         for j in range(num_train):
72             dists[i,j] = euclidean_dist(train_imgs[j,:], test_imgs[i,:])
73     return dists
74
75 def find_nearest_neighbors(k, dists, train_cls):
76     nn_distances = np.zeros((dists.shape[0],k))
77     nn_classes = np.zeros((dists.shape[0],k))
78     nn_classes_indices = np.zeros((dists.shape[0],k))
79     for i in range(dists.shape[0]):
80         idx = np.argpartition(dists[i,:], k, axis = 0)[:k]
81         nn_distances[i,:] = dists[i,idx]
82         nn_classes[i,:] = train_cls[idx]
83         nn_classes_indices[i,:] = idx
84     return nn_distances, nn_classes.astype(int), nn_classes_indices.astype(int)
85
86 def compute_error_rate(nn, test_cls, nn_cls):
87     acc_count = 0
88     total_test_imgs = nn.shape[0]
89     k = nn.shape[1]
90
91     for i in range(total_test_imgs):
92         unique_labels = np.unique(nn_cls[i])
93         label_count = 0
94         label = np.max(unique_labels) + 1
95         for j in range(unique_labels.shape[0]):
96             num_occurences = np.count_nonzero(nn_cls[i] == unique_labels[j])
97             if (num_occurences > label_count):
98                 label_count = num_occurences
99                 label = unique_labels[j]
100         elif(num_occurences == label_count):
101             new_indices = np.where(nn_cls[i] == unique_labels[j])[0]
102             old_indices = np.where(nn_cls[i] == label)[0]
103             new_avg = np.average(nn[i,new_indices])
104             old_avg = np.average(nn[i,old_indices])

```

```

105         if (new_avg < old_avg):
106             label = unique_labels[j]
107
108         if (test_cls[i] != label):
109             acc_count = acc_count + 1
110
111     return acc_count / total_test_imgs
112
113 def plot_nearest_neighbors(labels_to_keep, dists, train_cls, test_cls, test_imgs,
114                             train_imgs, class_names):
115     for i in range(labels_to_keep.shape[0]):
116         label = labels_to_keep[i]
117         nn_dists, nn_cls, nn_cls_indices = find_nearest_neighbors(5, dists, train_cls)
118         label_indices = np.where(test_cls == label)[0]
119         ind = np.random.randint(0, label_indices.shape[0] - 1)
120         label_index = label_indices[ind]
121         nn_img_indices = nn_cls_indices[label_index,:]
122         nn_cls = nn_cls[label_index]
123         nn_dists = nn_dists[label_index]
124
125         plt.figure(1)
126         ax1 = plt.subplot(2,3,1)
127         ax1.axis('off')
128         imgplt1 = plt.imshow(test_imgs[label_index])
129         ax1.set_title('Test Image ({}).format(class_names[i]), size=10)
130
131         for j in range(5):
132             ax2 = plt.subplot(2,3,(j+2))
133             ax2.axis('off')
134             nn_class = class_names[np.where(labels_to_keep == nn_cls[j])[0][0]]
135             imgplt2 = plt.imshow(train_imgs[nn_img_indices[j]])
136             ax2.set_title('NN: {} , dist: {:.2f}, class: {}'.format(j+1, nn_dists[j]
137                                     ,nn_class), size=7)
138
139         plt.show()
140
141 def plot_k_error_rates(k_array, dists, train_cls, test_cls):
142     e_rates = np.zeros(k_array.shape)
143     for k in range(k_array.shape[0]):
144         nn, nn_cls, _ = find_nearest_neighbors(k_array[k], dists, train_cls)
145         e_rate = compute_error_rate(nn, test_cls, nn_cls)
146         e_rates[k] = e_rate
147         print('Error rate for k = {} is {}'.format(k_array[k], e_rate))
148
149     plt.plot(k_array, e_rates)
150     plt.ylabel('Error Rates')
151     plt.xlabel('K')
152     plt.show()
153
154 def get_dists(filename, train_imgs, test_imgs):
155     # size of dists is 4000 x 10000
156     if (os.path.isfile(filename)):
157         dists = np.load(filename)

```

```

157     print('Precalculated distances loaded.')
158 else:
159     print('Calculating distances')
160     dists = compute_dists(train_imgs, test_imgs)
161     print('Saving distances')
162     np.save(filename, dists)
163     print('Distances saved')
164     return dists
165
166
167 def run_KNN():
168
169     # load CIFAR data
170     train_imgs, train_cls, train_names = cifar_loader.load_training_data()
171     test_imgs, test_cls, test_names = cifar_loader.load_test_data()
172     class_names = np.array(cifar_loader.load_class_names())
173
174     # get relevant class data
175     labels_to_keep = np.sort(np.array([0, 1, 8, 9]))
176     train_imgs, train_cls = split_sets(train_imgs, train_cls, labels_to_keep)
177     test_imgs, test_cls = split_sets(test_imgs, test_cls, labels_to_keep)
178     class_names = split_class_names(class_names, labels_to_keep)
179
180     # part a
181     view_examples(train_imgs, train_cls, labels_to_keep, class_names)
182
183     # part b and c
184     train_imgs_gray = rgb2gray_luminosity(train_imgs)
185     test_imgs_gray = rgb2gray_luminosity(test_imgs)
186     dists = get_dists('dists_lum_eucl.npy', train_imgs_gray, test_imgs_gray)
187
188
189
190     k_array = np.array([1,2,5,10,20])
191     plot_k_error_rates(k_array, dists, train_cls, test_cls)
192
193
194     # part d
195     plot_nearest_neighbors(labels_to_keep, dists, train_cls, test_cls, test_imgs,
196                             train_imgs, class_names)
197
198
199
200
201
202 if __name__ == '__main__':
203     run_KNN()

```

where the *cifar_loader* file is included in the code. The assignment is run through the function *run_KNN()*.

Part a: Random Image Visualization

I've chosen the classes

`'[airplane','automobile','ship','truck']`

which correspond with labels

`[0, 1, 8, 9]`

The random image visualization is accomplished by the function `view_examples()` and is called on line 181, with the function definition on lines 28-42. Calling the code once returns the figure below



Figure 1: Random Image Examples 1

Proving that it is indeed random, calling it again returns



Figure 2: Random Image Examples 2

Part b and c: Applying KNN Algorithm for $k = 1, 2, 5, 10, 20$

The relevant code in `run_KNN()` is in lines 184-191. This function converts all the training and testing images to grayscale to compute one distance between two images as opposed to one for each channel (lines 184-185). I used two different grayscale conversions, one being the average of the rgb channels, and the other being a coefficient based conversion. This last algorithm emphasizes that the majority of color that we see is in the green channel. The code then computes a distance matrix on line 186, with the function `get_dists()`, which looks for a saved distance matrix and loads it if it exists, and if not found, computes it. I the Euclidean Distance as a distance metric. The function `get_dists()` is on lines 153-164 and calls `compute_dists()`, which is on lines 65-73. Once the distances have been calculated or loaded, it computes the error rates and plots them using `plot_k_error_rates()` (lines 140-151), `find_nearest_neighbors()` (lines 75-84) and `compute_error_rate()` (lines 86-111). If the predicted classes for the nearest neighbors are the same, it chooses the neighbor with the smallest average distance. The results for the different grayscale conversions are below

Average Grayscale and Euclidean Distance

Using the average grayscale and Euclidean Distance gives the following output and graph

```
1 Error rate for k = 1 is 0.52975
2 Error rate for k = 2 is 0.52975
3 Error rate for k = 5 is 0.53325
4 Error rate for k = 10 is 0.5355
5 Error rate for k = 20 is 0.549
```




Figure 3: K vs Error Rate Using Average Grayscale and Euclidean Distance

A more comprehensive sweep, using k's from 1-20 gives

```
1 Error rate for k = 1 is 0.52975
2 Error rate for k = 2 is 0.52975
3 Error rate for k = 3 is 0.52925
4 Error rate for k = 4 is 0.53025
5 Error rate for k = 5 is 0.53325
6 Error rate for k = 6 is 0.5305
7 Error rate for k = 7 is 0.5355
8 Error rate for k = 8 is 0.53675
9 Error rate for k = 9 is 0.53325
10 Error rate for k = 10 is 0.5355
```



```

11 Error rate for k = 11 is 0.537
12 Error rate for k = 12 is 0.53825
13 Error rate for k = 13 is 0.5405
14 Error rate for k = 14 is 0.54175
15 Error rate for k = 15 is 0.54175
16 Error rate for k = 16 is 0.54525
17 Error rate for k = 17 is 0.54825
18 Error rate for k = 18 is 0.54875
19 Error rate for k = 19 is 0.54825
20 Error rate for k = 20 is 0.549

```




Figure 4: K vs Error Rate Using Average Grayscaleing and Euclidean Distance

Which shows that the minimum error rate is again given by $k = 3$, and the error rate for $k = 1$ is 0.52975

Coefficient Grayscale and Euclidean Distance

Using the coefficient grayscale and Euclidean Distance gives the following output and graph

```
1 Error rate for k = 1 is 0.53125
2 Error rate for k = 2 is 0.53125
3 Error rate for k = 5 is 0.54
4 Error rate for k = 10 is 0.54275
5 Error rate for k = 20 is 0.5475
```




Figure 5: K vs Error Rate Using Coefficient Grayscale and Euclidean Distance

A more comprehensive sweep, using k's from 1-20 gives

```
1 Error rate for k = 1 is 0.53125
2 Error rate for k = 2 is 0.53125
3 Error rate for k = 3 is 0.53875
4 Error rate for k = 4 is 0.54025
5 Error rate for k = 5 is 0.54
6 Error rate for k = 6 is 0.542
7 Error rate for k = 7 is 0.545
8 Error rate for k = 8 is 0.54075
9 Error rate for k = 9 is 0.54125
10 Error rate for k = 10 is 0.54275
```

```

11 Error rate for k = 11 is 0.54075
12 Error rate for k = 12 is 0.546
13 Error rate for k = 13 is 0.54575
14 Error rate for k = 14 is 0.54625
15 Error rate for k = 15 is 0.546
16 Error rate for k = 16 is 0.5465
17 Error rate for k = 17 is 0.54625
18 Error rate for k = 18 is 0.5485
19 Error rate for k = 19 is 0.54475
20 Error rate for k = 20 is 0.5475

```




Figure 6: K vs Error Rate Using Coefficient Grayscale and Euclidean Distance

Which shows that the minimum error rate is again given by $k = 1, 2$, and the error rate for $k = 1$ is 0.53125.

In general, the error rate does not decrease with k , and it shouldn't since it is a dataset dependent variable.

1 Part d: Visualizing the 5 Nearest Neighbors

This part is accomplished with the function call on 196 (definition on lines 113-138). Two runs of the code gives the plots below which show the randomness of the image choice.

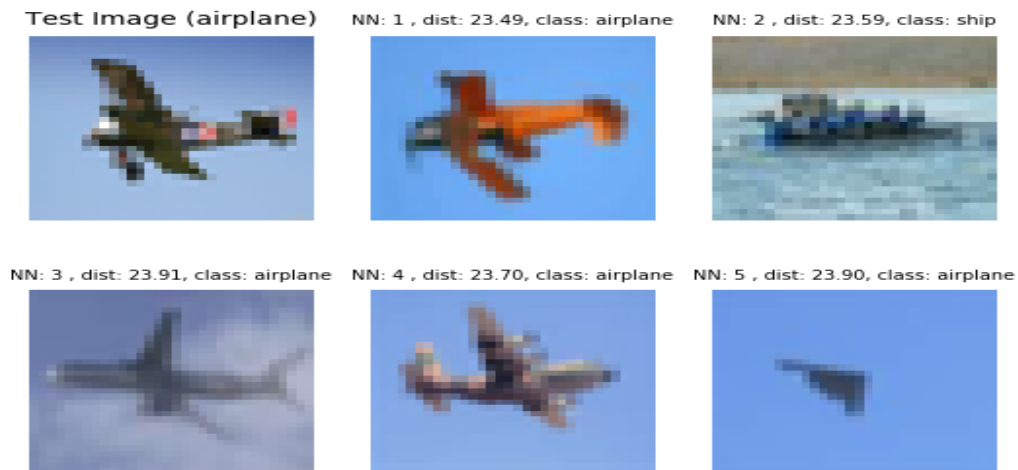


Figure 7: Random Airplane Image and it's 5 Nearest Neighbors



Figure 8: Random Airplane Image and it's 5 Nearest Neighbors



Figure 9: Random Automobile Image and it's 5 Nearest Neighbors



Figure 10: Random Automobile Image and it's 5 Nearest Neighbors



Figure 11: Random Ship Image and it's 5 Nearest Neighbors

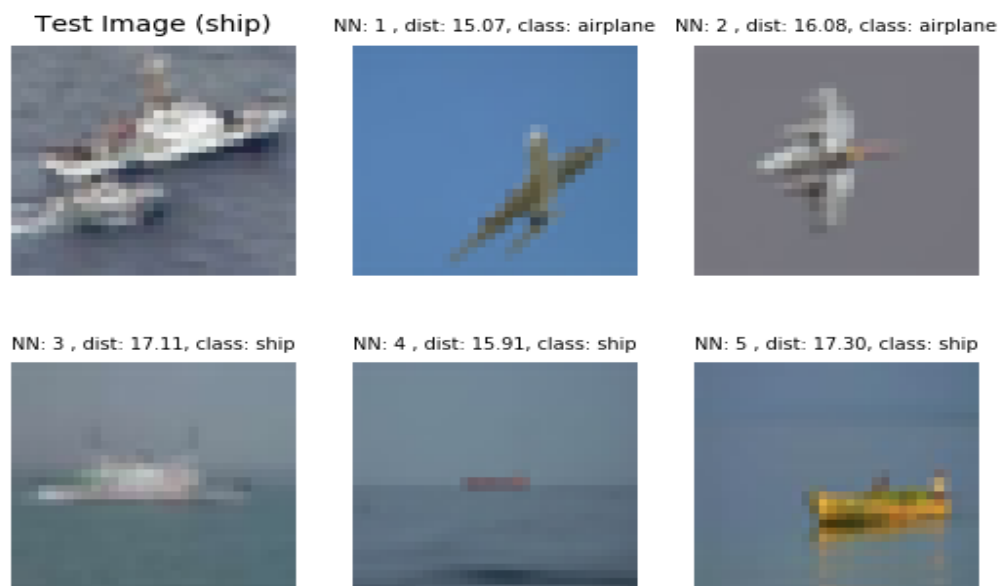


Figure 12: Random Ship Image and it's 5 Nearest Neighbors

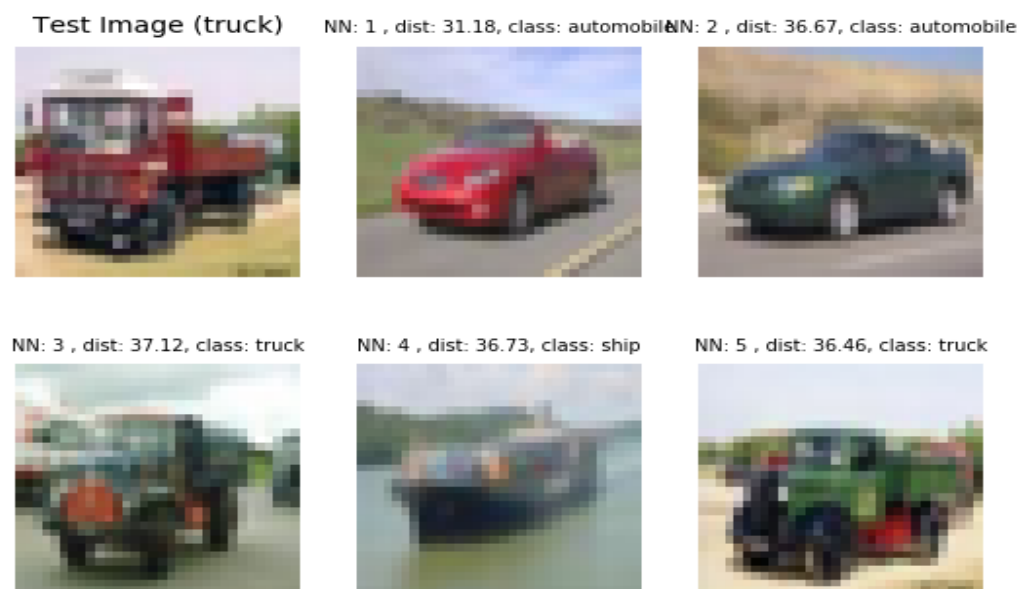


Figure 13: Random Truck Image and it's 5 Nearest Neighbors




Figure 14: Random Truck Image and it's 5 Nearest Neighbors