

ECE194N HW 2

CNN Report

February 22, 2019

Student: Erik Rosten
Perm Number: 7143571
Email: erosten@ucsb.edu

Department of Electrical and Computer Engineering, UCSB

Part a: Visualizing the Dataset

The code for this section is in `view_examples.py` and repeated below

```
1 import cifar_loader
2 import matplotlib.pyplot as plt
3 import numpy as np
4
5
6 def view_examples(train_imgs, train_cls, classes, class_names):
7     print(class_names)
8     plt.figure(1)
9     for i in range(classes.shape[0]):
10         label_class_name = class_names[i]
11         rand_integer = np.random.randint(0,4999)
12         # grab a random example index
13         index = np.where(train_cls == classes[i])[0][rand_integer]
14         img_example = train_imgs[index]
15
16         ax = plt.subplot(2,classes.shape[0] / 2,i + 1)
17         ax.axis('off')
18         imgplot = plt.imshow(img_example)
19         ax.set_title(label_class_name)
20     plt.show()
21
22
23 def main():
24     # load cifar data
25     train_imgs, train_cls, train_names = cifar_loader.load_training_data()
26     class_names = np.array(cifar_loader.load_class_names())
27     classes = np.arange(0,10)
28     # view examples
29     view_examples(train_imgs, train_cls, classes, class_names)
30
31
32
33
34 if __name__ == '__main__':
35     main()
```

Running the main function of this code gives the figure below



Figure 1: One example from each class in the cifar-10 dataset

Designing the CNN and Results

The code for this section is in `train_cnn.py` and repeated below

```

1 from __future__ import print_function
2 import keras
3 from keras.datasets import cifar10
4 from keras.preprocessing.image import ImageDataGenerator
5 from keras.models import Sequential
6 from keras.layers import Dense, Activation, Flatten, Dropout
7 from keras.layers import Conv2D, MaxPooling2D
8 from keras.layers import LeakyReLU
9 from keras.layers.normalization import BatchNormalization
10 import os
11 import numpy as np
12
13 # define some variables
14 batch_size = 32
15 num_classes = 10
16 epochs = 25
17 data_augmentation = True

```

```

18 num_predictions = 20
19 save_dir = os.path.join(os.getcwd(), 'saved_models')
20 model_name = 'keras_cifar10_modified_vgg_dropout.h5'
21
22 # load cifar data
23 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
24 print('x_train shape:', x_train.shape)
25 print(x_train.shape[0], 'train samples')
26 print(x_test.shape[0], 'test samples')
27
28 # Convert class vectors to one hot
29 y_train = keras.utils.to_categorical(y_train, num_classes)
30 y_test = keras.utils.to_categorical(y_test, num_classes)
31
32 # create modified VGG
33 model = Sequential()
34 #block 1
35 model.add(Conv2D(64, (3, 3), padding='same',
36                 input_shape=x_train.shape[1:]))
37 model.add(LeakyReLU(alpha = 0.1))
38 model.add(BatchNormalization())
39 model.add(Conv2D(64, (3, 3), padding='same'))
40 model.add(LeakyReLU(alpha = 0.1))
41 model.add(BatchNormalization())
42 model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
43
44 #block 2
45 model.add(Conv2D(128, (3, 3), padding='same'))
46 model.add(BatchNormalization())
47 model.add(LeakyReLU(alpha = 0.2))
48 model.add(Conv2D(128, (3, 3), padding='same'))
49 model.add(LeakyReLU(alpha = 0.2))
50 model.add(BatchNormalization())
51 model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
52
53
54 #block 3
55 model.add(Conv2D(256, (3, 3), padding='same'))
56 model.add(LeakyReLU(alpha = 0.2))
57 model.add(BatchNormalization())
58 model.add(Conv2D(256, (3, 3), padding='same'))
59 model.add(LeakyReLU(alpha = 0.2))
60 model.add(BatchNormalization())
61 model.add(Conv2D(256, (3, 3), padding='same'))
62 model.add(LeakyReLU(alpha = 0.2))
63 model.add(BatchNormalization())
64 model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
65
66
67 model.add(Flatten())
68 model.add(Dense(4096))
69 model.add(LeakyReLU(alpha = 0.2))
70 model.add(Dropout(0.5))
71 model.add(Dense(4096))

```

```

72 model.add(LeakyReLU(alpha = 0.2))
73 model.add(Dropout(0.5))
74 model.add(Dense(num_classes))
75 model.add(LeakyReLU(alpha = 0.2))
76 model.add(Activation('softmax'))
77
78 # initiate Adam optimizer
79 opt = keras.optimizers.Adam(epsilon = 1e-03)
80
81 model.compile(loss='categorical_crossentropy',
82               optimizer=opt,
83               metrics=['accuracy'])
84
85 x_train = x_train.astype('float32')
86 x_test = x_test.astype('float32')
87 x_train /= 255
88 x_test /= 255
89
90
91 # This will do preprocessing and realtime data augmentation:
92 datagen = ImageDataGenerator(
93     featurewise_center=False, # set input mean to 0 over the dataset
94     samplewise_center=False, # set each sample mean to 0
95     featurewise_std_normalization=False, # divide inputs by std of the dataset
96     samplewise_std_normalization=False, # divide each input by its std
97     zca_whitening=False, # apply ZCA whitening
98     zca_epsilon=1e-06, # epsilon for ZCA whitening
99     rotation_range=0, # randomly rotate images in the range (degrees, 0 to 180)
100     # randomly shift images horizontally (fraction of total width)
101     width_shift_range=0.1,
102     # randomly shift images vertically (fraction of total height)
103     height_shift_range=0.1,
104     shear_range=0., # set range for random shear
105     zoom_range=0., # set range for random zoom
106     channel_shift_range=0., # set range for random channel shifts
107     # set mode for filling points outside the input boundaries
108     fill_mode='nearest',
109     cval=0., # value used for fill_mode = "constant"
110     horizontal_flip=True, # randomly flip images
111     vertical_flip=False, # randomly flip images
112     # set rescaling factor (applied before any other transformation)
113     rescale=None,
114     # set function that will be applied on each input
115     preprocessing_function=None,
116     # image data format, either "channels_first" or "channels_last"
117     data_format=None,
118     # fraction of images reserved for validation (strictly between 0 and 1)
119     validation_split=0.0)
120
121 # Compute quantities required for feature-wise normalization
122 # (std, mean, and principal components if ZCA whitening is applied).
123 datagen.fit(x_train)
124
125 # Fit the model on the batches generated by datagen.flow().

```

```

126 history = model.fit_generator(datagen.flow(x_train, y_train,
127                                     batch_size=batch_size),
128                             steps_per_epoch = x_train.shape[0] / batch_size,
129                             epochs=epochs,
130                             validation_data=(x_test, y_test),
131                             workers=4, verbose = 1)
132 train_loss = history.history['loss']
133 test_loss = history.history['val_loss']
134 train_acc = history.history['acc']
135 test_acc = history.history['val_acc']
136 np.savetxt("train_loss.txt", train_loss, delimiter=",")
137 np.savetxt("train_acc.txt", train_acc, delimiter=",")
138 np.savetxt("test_loss.txt", test_loss, delimiter=",")
139 np.savetxt("test_acc.txt", test_acc, delimiter=",")
140
141
142 # Save model and weights
143 if not os.path.isdir(save_dir):
144     os.makedirs(save_dir)
145 model_path = os.path.join(save_dir, model_name)
146 model.save(model_path)
147 print('Saved trained model at %s ' % model_path)
148
149 # Score trained model.
150 scores = model.evaluate(x_test, y_test, verbose=1)
151 print('Test loss:', scores[0])
152 print('Test accuracy:', scores[1])

```

For this section, I chose a modified VGG16. I've taken the VGG16, removed the last 6 conv/pooling layers, and added dropout to the second two fully connected layers. It is therefore 10 layers, consisting of 7 convolutional layers, and 3 fully connected layers. The loss function is cross entropy, which is used to find the loss of two different probability distributions. This assumes that the data has an underlying probability distribution that the CNN is trying to match. The code above saves the train and test losses in text files after 25 epochs of training, which have been included in this zip file. The plots for training and test accuracy and loss is below, and calculated using the code below in plot_results.py, repeated here

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4
5 test_loss = np.loadtxt('test_loss.txt')
6 test_acc = np.loadtxt('test_acc.txt')
7 train_loss = np.loadtxt('train_loss.txt')
8 train_acc = np.loadtxt('train_acc.txt')
9
10
11 epochs = np.arange(test_loss.shape[0])
12
13 fig = plt.figure(1)
14 # plot losses
15 ax = plt.subplot(1,2,1)
16 plt.plot(epochs, train_loss)

```

```

17 plt.plot(epochs, test_loss)
18 ax.legend(['Training', 'Test'])
19 ax.set_title('Training and Test Loss vs Epoch')
20 ax.set_xlabel('Epoch')
21 # plot accuracies
22 ax = plt.subplot(1,2,2)
23 plt.plot(epochs, train_acc)
24 plt.plot(epochs, test_acc)
25 ax.legend(['Training', 'Test'])
26 ax.set_title('Training and Test Accuracy vs Epoch')
27 ax.set_xlabel('Epoch')
28
29
30 fig.tight_layout()
31 plt.show()

```

which gives the plot below

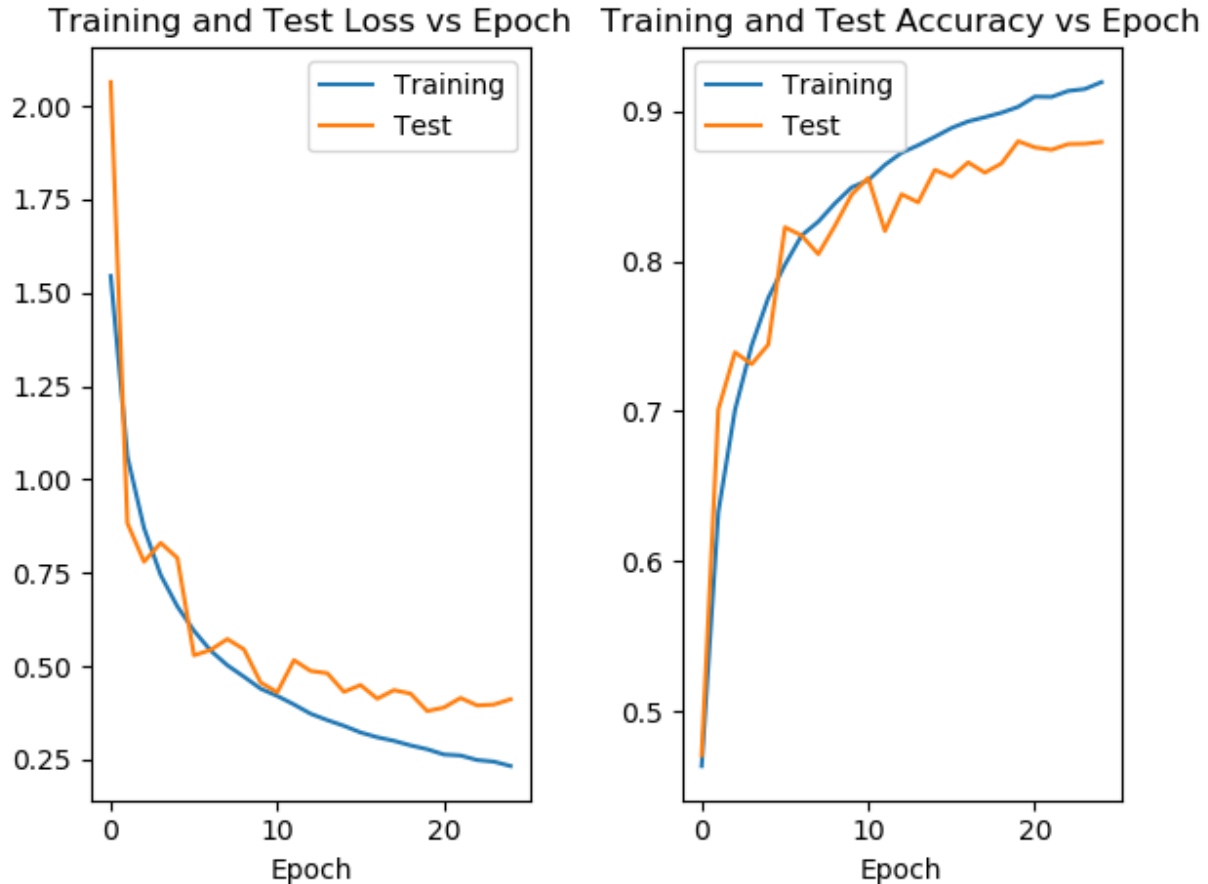


Figure 2: Results of the Modified VGG

Note that there is a good amount of overfitting after around epoch 10, which is what drove my choice to include dropout in the network. This is another reason for using real time data augmentation. The accuracy could be improved by perhaps adding more techniques that combat overfitting.

Part c: Model Parameters

One advantage of using keras is it's model summary function. Using the code below from `print_num_parameters.py`

```
1 from __future__ import print_function
2 import keras
3 from keras.models import load_model
4 from keras.datasets import cifar10
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8
9 model = load_model('saved_models/keras_cifar10_modified_vgg_dropout.h5')
10 model.summary()
```

gives

```
1 -----
2 Layer (type)                Output Shape          Param #
3 -----
4 conv2d_1 (Conv2D)           (None, 32, 32, 64)    1792
5 -----
6 leaky_re_lu_1 (LeakyReLU)   (None, 32, 32, 64)    0
7 -----
8 batch_normalization_1 (Batch Normalization) (None, 32, 32, 64)    256
9 -----
10 conv2d_2 (Conv2D)           (None, 32, 32, 64)    36928
11 -----
12 leaky_re_lu_2 (LeakyReLU)   (None, 32, 32, 64)    0
13 -----
14 batch_normalization_2 (Batch Normalization) (None, 32, 32, 64)    256
15 -----
16 max_pooling2d_1 (MaxPooling2D) (None, 16, 16, 64)    0
17 -----
18 conv2d_3 (Conv2D)           (None, 16, 16, 128)   73856
19 -----
20 batch_normalization_3 (Batch Normalization) (None, 16, 16, 128)   512
21 -----
22 leaky_re_lu_3 (LeakyReLU)   (None, 16, 16, 128)   0
23 -----
24 conv2d_4 (Conv2D)           (None, 16, 16, 128)   147584
25 -----
26 leaky_re_lu_4 (LeakyReLU)   (None, 16, 16, 128)   0
27 -----
28 batch_normalization_4 (Batch Normalization) (None, 16, 16, 128)   512
29 -----
30 max_pooling2d_2 (MaxPooling2D) (None, 8, 8, 128)    0
31 -----
32 conv2d_5 (Conv2D)           (None, 8, 8, 256)     295168
33 -----
34 leaky_re_lu_5 (LeakyReLU)   (None, 8, 8, 256)     0
35 -----
36 batch_normalization_5 (Batch Normalization) (None, 8, 8, 256)     1024
37 -----
38 conv2d_6 (Conv2D)           (None, 8, 8, 256)     590080
```



```

39 -----
40 leaky_re_lu_6 (LeakyReLU) (None, 8, 8, 256)      0
41 -----
42 batch_normalization_6 (Batch Normalization) (None, 8, 8, 256) 1024
43 -----
44 conv2d_7 (Conv2D) (None, 8, 8, 256)      590080
45 -----
46 leaky_re_lu_7 (LeakyReLU) (None, 8, 8, 256)      0
47 -----
48 batch_normalization_7 (Batch Normalization) (None, 8, 8, 256) 1024
49 -----
50 max_pooling2d_3 (MaxPooling2D) (None, 4, 4, 256) 0
51 -----
52 flatten_1 (Flatten) (None, 4096)      0
53 -----
54 dense_1 (Dense) (None, 4096)      16781312
55 -----
56 leaky_re_lu_8 (LeakyReLU) (None, 4096)      0
57 -----
58 dropout_1 (Dropout) (None, 4096)      0
59 -----
60 dense_2 (Dense) (None, 4096)      16781312
61 -----
62 leaky_re_lu_9 (LeakyReLU) (None, 4096)      0
63 -----
64 dropout_2 (Dropout) (None, 4096)      0
65 -----
66 dense_3 (Dense) (None, 10)      40970
67 -----
68 leaky_re_lu_10 (LeakyReLU) (None, 10)      0
69 -----
70 activation_1 (Activation) (None, 10)      0
71 =====
72 Total params: 35,343,690
73 Trainable params: 35,341,386
74 Non-trainable params: 2,304
75 -----

```

Part d: Viewing Wrongly Classified Images

This section utilized the code from `find_incorrect_instances.py`.

```

1 from __future__ import print_function
2 import keras
3 from keras.models import load_model
4 from keras.datasets import cifar10
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8
9 #load model
10 model = load_model('saved_models/keras_cifar10_modified_vgg_dropout.h5')
11 print('Model Loaded')

```

```

12
13
14 # load cifar data
15 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
16 print('Dataset Loaded')
17
18
19 # preprocess
20 x_train = x_train.astype('float32')
21 x_test = x_test.astype('float32')
22 x_train /= 255
23 x_test /= 255
24
25
26 # find incorrectly classified images
27 print('Finding Incorrectly Classified Images')
28 predictions = model.predict_classes(x_test).reshape((-1,1))
29 incorrects = np.nonzero(predictions!= y_test)[0]
30 np.random.shuffle(incorrects)
31 _, incorrect_indices = np.unique(y_test[incorrects], return_index = True)
32 incorrect_indices = incorrects[incorrect_indices]
33 for i in range(incorrect_indices.shape[0]):
34     index = incorrect_indices[i]
35     imgplt = plt.imshow(x_test[index])
36     plt.title('Predicted Class: {} Actual Class: {}'.format(predictions[index],
37         y_test[index]))
37     plt.show()

```

A reminder of the classes in CIFAR-10 from 0 to 10 is below

[airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck]

Running the code above gives the incorrect images below

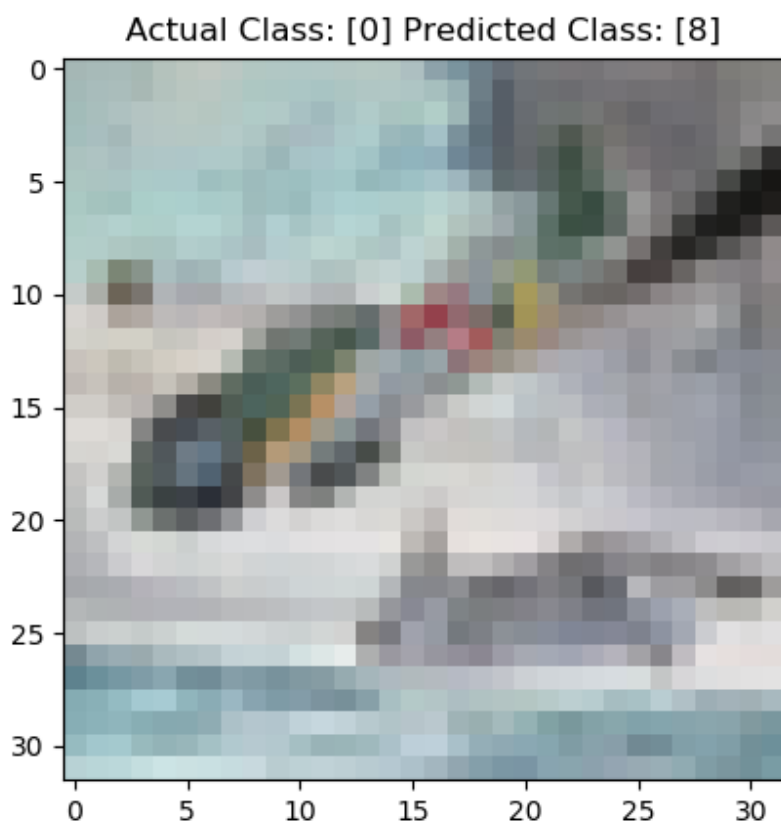


Figure 3: Real: airplane, Predicted: ship

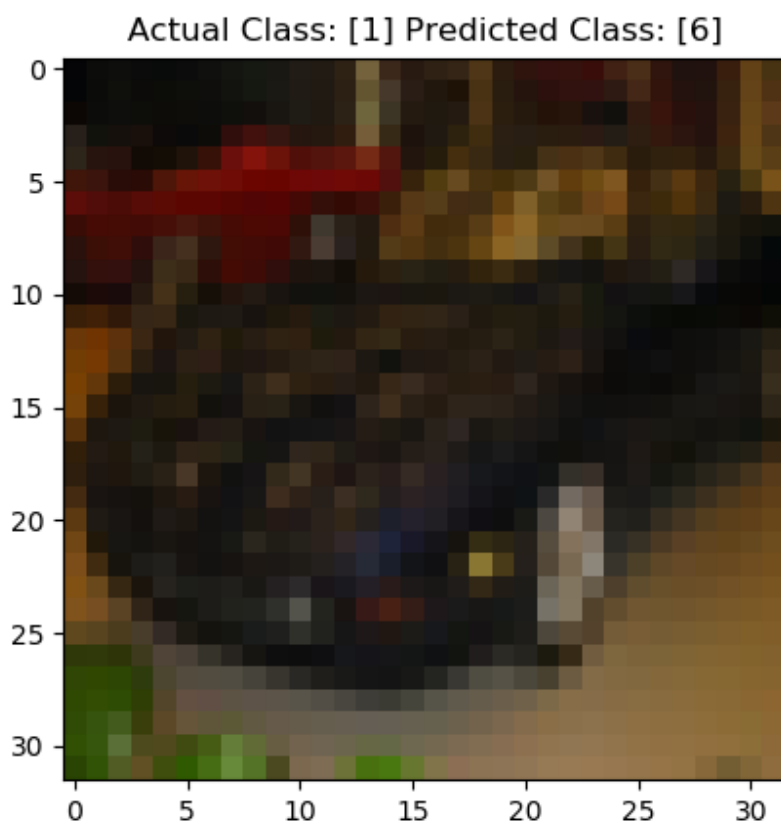


Figure 4: Real: automobile, Predicted: frog

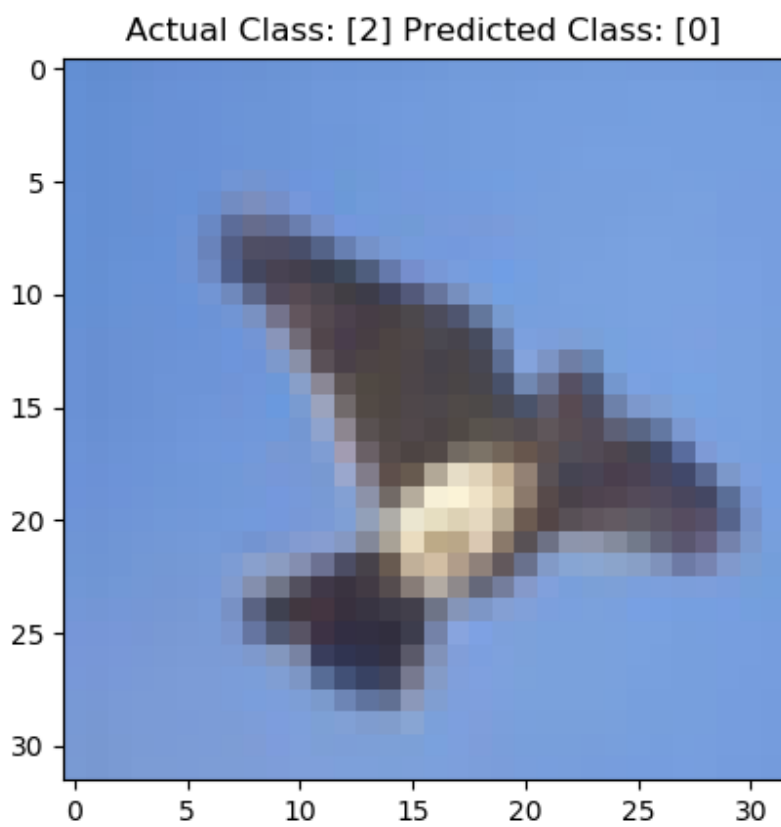


Figure 5: Real: bird, Predicted: airplane

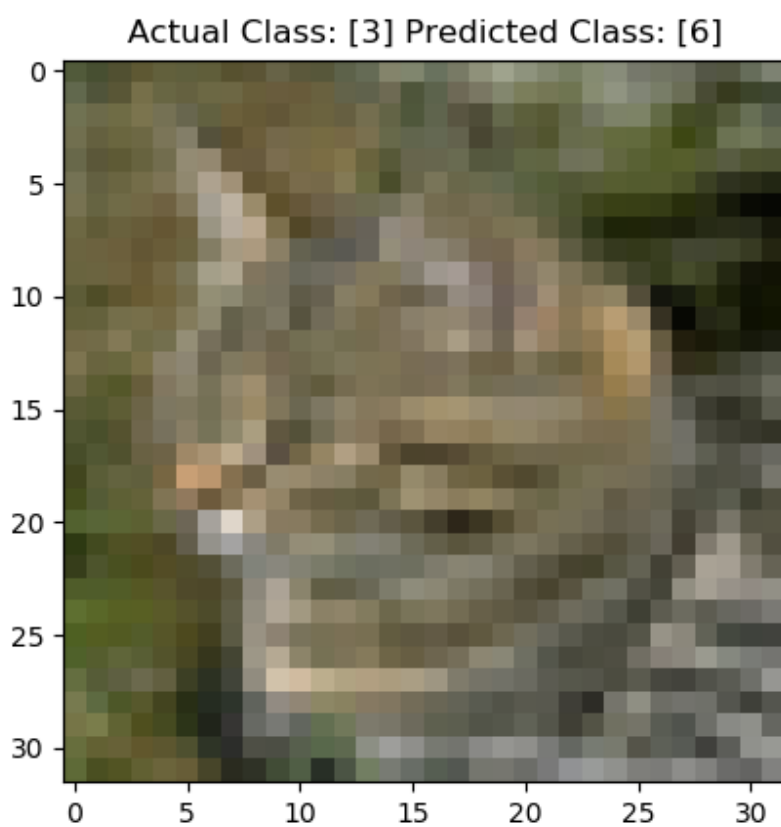


Figure 6: Real: cat, Predicted: frog

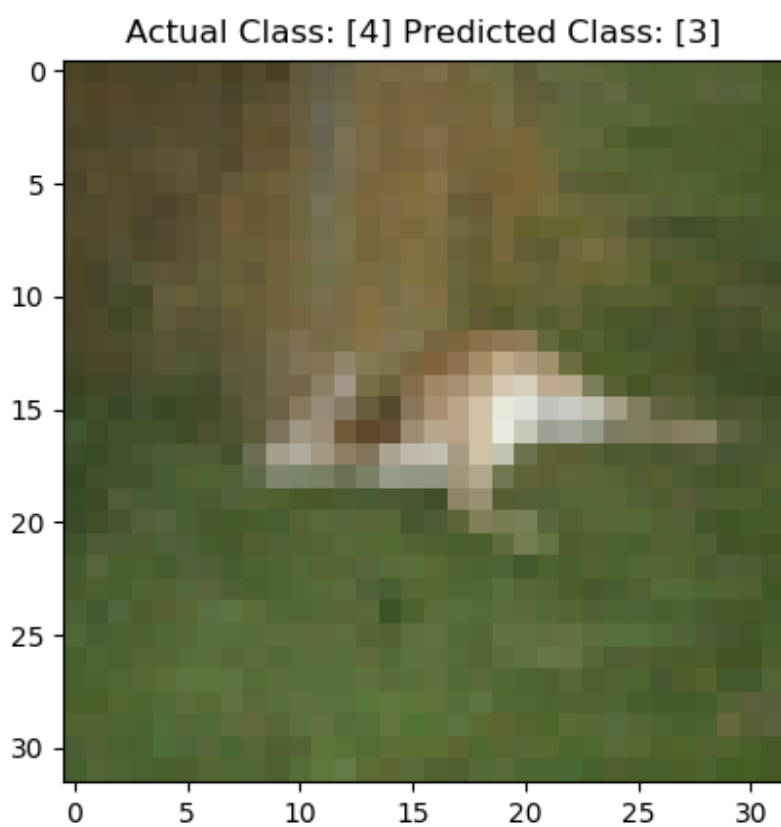


Figure 7: Real: deer, Predicted: cat

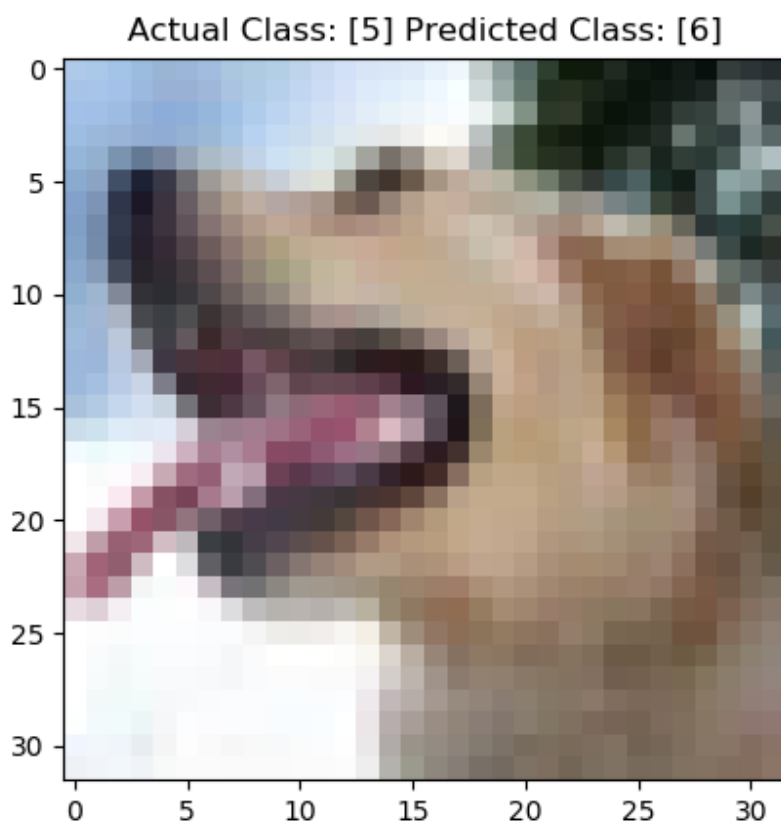


Figure 8: Real: dog, Predicted: frog




Figure 9: Real: frog, Predicted: automobile

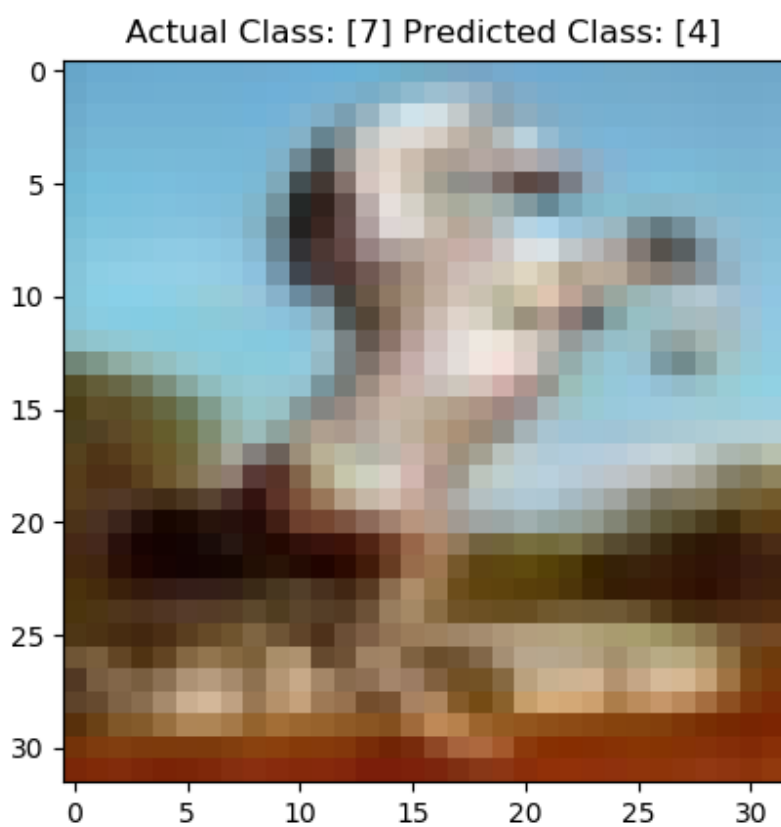


Figure 10: Real: horse, Predicted: deer

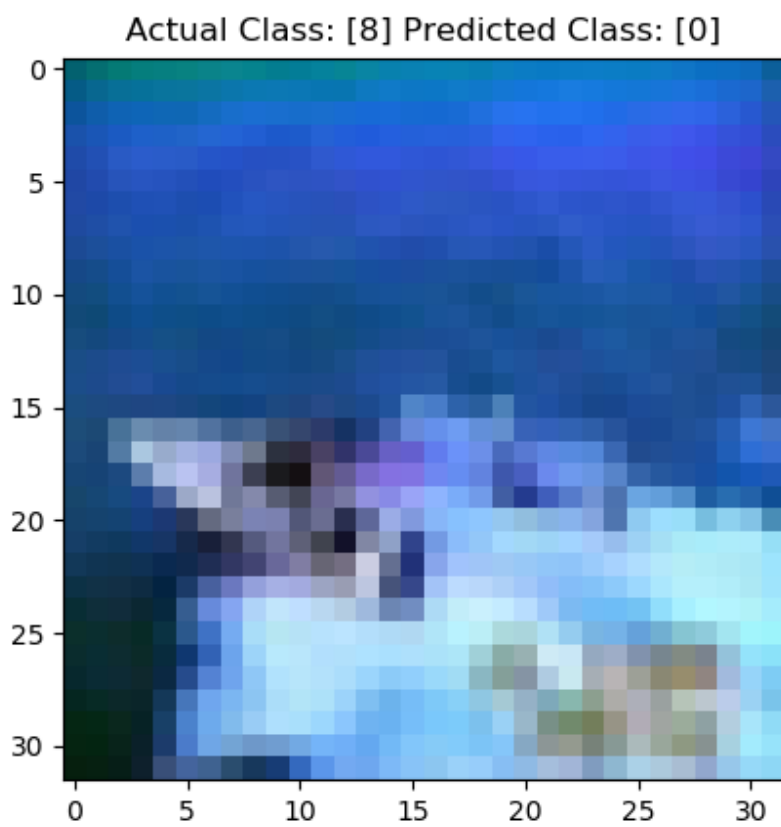


Figure 11: Real: ship, Predicted: airplane




Figure 12: Real: truck, Predicted: deer