

# ECE194N HW 3

---

## RNN Report

March 8, 2019

**Student:** Erik Rosten  
**Perm Number:** 7143571  
**Email:** erosten@ucsb.edu

**Department of Electrical and Computer Engineering, UCSB**

## a. Visualizing the Data

The code for visualizing the data is below

```
1 import reader
2 import numpy as np
3
4
5 raw_data = reader.ptb_raw_data('simple-examples/data/')
6 train_data, valid_data, test_data, _, word_dict = raw_data
7
8 train_data = np.array(train_data)
9 valid_data = np.array(valid_data)
10 test_data = np.array(test_data)
11 word_dict = dict(zip(word_dict.values(),word_dict.keys()))
12
13
14
15 sentence = ''
16 keys = []
17 sentenceNum = 1
18 for j in range(valid_data.shape[0]):
19     ind = valid_data[j]
20     if (word_dict[ind] == '<eos>'):
21         print('\n\n')
22         print('Sentence {}'.format(sentenceNum))
23         print('\{"{}\\" {} '.format(sentence, np.array(keys)))
24         sentence = ''
25         keys = []
26         sentenceNum = sentenceNum + 1
27         if (sentenceNum == 11):
28             break
29         continue
30
31     if (j > 0):
32         sentence = sentence + ' '
33
34     sentence = sentence + word_dict[ind]
35     keys.append(ind)
```

where I have modified the reader code to return the word\_dict. Running this code returns

```
1
2 Sentence 1
3 "consumers may want to move their telephones a little closer to the tv set" [1132 93 358
   5 329 51 9836 6 326 2476 5 0 662 388]
4
5
6
7 Sentence 2
8 " <unk> <unk> watching abc 's monday night football can now vote during <unk> for the
   greatest play in N years from among four or five <unk> <unk>" [ 1 1 2974 2158 9 381
   1068 2347 89 99 847 198 1 11
9 0 3383 1119 7 3 72 20 211 346 36 258 1 1]
10
```

11  
12  
13 Sentence 3  
14 " two weeks ago viewers of several nbc <unk> consumer segments started calling a N  
number for advice on various <unk> issues" [ 75 422 195 3917 4 249 1795 1 580 3528  
892 2374 6 3  
15 297 11 2709 16 1186 1 250]  
16  
17  
18  
19 Sentence 4  
20 " and the new syndicated reality show hard copy records viewers ' opinions for possible  
airing on the next day 's show" [ 8 0 35 9922 3747 464 710 2998 2037 3917 134 6145  
11 494  
21 5894 16 0 130 272 9 464]  
22  
23  
24  
25 Sentence 5  
26 " interactive telephone technology has taken a new leap in <unk> and television  
programmers are racing to exploit the possibilities" [9958 732 503 30 641 6 35 6498  
7 1 8 761 9967 26  
27 6587 5 6415 0 6574]  
28  
29  
30  
31 Sentence 6  
32 " eventually viewers may grow <unk> with the technology and <unk> the cost" [1413 3917  
93 1552 1 22 0 503 8 1 0 361]  
33  
34  
35  
36 Sentence 7  
37 " but right now programmers are figuring that viewers who are busy dialing up a range of  
services may put down their <unk> control <unk> and stay <unk>" [ 29 382 99 9967 26  
7428 10 3917 56 26 3248 8846 52 6  
38 880 4 323 93 335 118 51 1 350 1 8 1337 1]  
39  
40  
41  
42 Sentence 8  
43 " we 've been spending a lot of time in los angeles talking to tv production people says  
mike parks president of call interactive which supplied technology for both abc  
sports and nbc 's consumer minutes" [ 64 573 58 508 6 581 4 103 7 639 747 1921 5 662  
44 359 108 44 5458 6149 70 4 786 9958 41 7746 503 11 179  
45 2158 1259 8 1795 9 580 1495]  
46  
47  
48  
49 Sentence 9  
50 " with the competitiveness of the television market these days everyone is looking for a  
way to get viewers more excited" [ 22 0 9643 4 0 761 47 144 171 1376 13 735 11 6  
51 229 5 188 3917 45 9684]  
52

```

53
54
55 Sentence 10
56 " one of the leaders behind the expanded use of N numbers is call interactive a joint
    venture of giants american express co. and american telephone & telegraph co" [ 54 4
    0 815 1116 0 2439 269 4 3 1619 13 786 9958
57 6 795 818 4 2172 140 1021 95 8 140 732 82 3133 570]

```

## b. Training the model

Modifying the given code to save loss and running the code gives

```

1 Epoch: 1 Learning rate: 1.000
2 0.004 perplexity: 6695.571 speed: 8623 wps
3 0.104 perplexity: 841.522 speed: 21685 wps
4 0.204 perplexity: 618.627 speed: 22346 wps
5 0.304 perplexity: 500.656 speed: 22739 wps
6 0.404 perplexity: 432.179 speed: 22809 wps
7 0.504 perplexity: 387.865 speed: 22326 wps
8 0.604 perplexity: 349.699 speed: 22334 wps
9 0.703 perplexity: 323.348 speed: 22494 wps
10 0.803 perplexity: 302.308 speed: 22624 wps
11 0.903 perplexity: 282.956 speed: 22716 wps
12 Epoch: 1 Train Perplexity: 268.422
13 Epoch: 1 Valid Perplexity: 179.301
14 Epoch: 2 Learning rate: 1.000
15 0.004 perplexity: 208.967 speed: 24266 wps
16 0.104 perplexity: 149.908 speed: 21422 wps
17 0.204 perplexity: 157.417 speed: 22361 wps
18 0.304 perplexity: 152.346 speed: 22704 wps
19 0.404 perplexity: 149.530 speed: 22890 wps
20 0.504 perplexity: 147.265 speed: 23018 wps
21 0.604 perplexity: 142.691 speed: 23108 wps
22 0.703 perplexity: 140.512 speed: 23173 wps
23 0.803 perplexity: 138.554 speed: 23199 wps
24 0.903 perplexity: 134.969 speed: 23223 wps
25 Epoch: 2 Train Perplexity: 132.864
26 Epoch: 2 Valid Perplexity: 142.850
27 Epoch: 3 Learning rate: 1.000
28 0.004 perplexity: 143.674 speed: 23878 wps
29 0.104 perplexity: 104.587 speed: 23382 wps
30 0.204 perplexity: 113.822 speed: 23413 wps
31 0.304 perplexity: 111.091 speed: 23445 wps
32 0.404 perplexity: 110.110 speed: 23460 wps
33 0.504 perplexity: 109.426 speed: 23489 wps
34 0.604 perplexity: 106.796 speed: 23493 wps
35 0.703 perplexity: 106.130 speed: 23490 wps
36 0.803 perplexity: 105.512 speed: 23481 wps
37 0.903 perplexity: 103.273 speed: 23481 wps
38 Epoch: 3 Train Perplexity: 102.256
39 Epoch: 3 Valid Perplexity: 132.594
40 Epoch: 4 Learning rate: 1.000
41 0.004 perplexity: 117.351 speed: 22860 wps

```

42 0.104 perplexity: 85.203 speed: 23450 wps  
 43 0.204 perplexity: 93.827 speed: 23464 wps  
 44 0.304 perplexity: 91.566 speed: 23470 wps  
 45 0.404 perplexity: 91.025 speed: 23452 wps  
 46 0.504 perplexity: 90.677 speed: 23475 wps  
 47 0.604 perplexity: 88.782 speed: 23485 wps  
 48 0.703 perplexity: 88.584 speed: 23494 wps  
 49 0.803 perplexity: 88.384 speed: 23501 wps  
 50 0.903 perplexity: 86.737 speed: 23512 wps  
 51 Epoch: 4 Train Perplexity: 86.164  
 52 Epoch: 4 Valid Perplexity: 128.316  
 53 Epoch: 5 Learning rate: 0.500  
 54 0.004 perplexity: 101.607 speed: 23739 wps  
 55 0.104 perplexity: 71.511 speed: 23541 wps  
 56 0.204 perplexity: 77.490 speed: 23531 wps  
 57 0.304 perplexity: 74.558 speed: 23525 wps  
 58 0.404 perplexity: 73.472 speed: 23517 wps  
 59 0.504 perplexity: 72.606 speed: 23505 wps  
 60 0.604 perplexity: 70.490 speed: 23519 wps  
 61 0.703 perplexity: 69.718 speed: 23540 wps  
 62 0.803 perplexity: 68.935 speed: 23554 wps  
 63 0.903 perplexity: 67.016 speed: 23546 wps  
 64 Epoch: 5 Train Perplexity: 65.996  
 65 Epoch: 5 Valid Perplexity: 119.056  
 66 Epoch: 6 Learning rate: 0.250  
 67 0.004 perplexity: 82.694 speed: 23130 wps  
 68 0.104 perplexity: 59.016 speed: 23488 wps  
 69 0.204 perplexity: 64.249 speed: 23440 wps  
 70 0.304 perplexity: 61.726 speed: 23426 wps  
 71 0.404 perplexity: 60.722 speed: 23531 wps  
 72 0.504 perplexity: 59.913 speed: 23481 wps  
 73 0.604 perplexity: 58.063 speed: 23460 wps  
 74 0.703 perplexity: 57.312 speed: 23466 wps  
 75 0.803 perplexity: 56.502 speed: 23462 wps  
 76 0.903 perplexity: 54.753 speed: 23464 wps  
 77 Epoch: 6 Train Perplexity: 53.782  
 78 Epoch: 6 Valid Perplexity: 118.198  
 79 Epoch: 7 Learning rate: 0.125  
 80 0.004 perplexity: 72.797 speed: 24196 wps  
 81 0.104 perplexity: 52.220 speed: 23343 wps  
 82 0.204 perplexity: 57.053 speed: 23377 wps  
 83 0.304 perplexity: 54.856 speed: 23410 wps  
 84 0.404 perplexity: 53.933 speed: 23446 wps  
 85 0.504 perplexity: 53.173 speed: 23493 wps  
 86 0.604 perplexity: 51.488 speed: 23498 wps  
 87 0.703 perplexity: 50.781 speed: 23494 wps  
 88 0.803 perplexity: 49.991 speed: 23476 wps  
 89 0.903 perplexity: 48.373 speed: 23483 wps  
 90 Epoch: 7 Train Perplexity: 47.455  
 91 Epoch: 7 Valid Perplexity: 119.343  
 92 Epoch: 8 Learning rate: 0.062  
 93 0.004 perplexity: 68.244 speed: 23580 wps  
 94 0.104 perplexity: 48.842 speed: 23515 wps  
 95 0.204 perplexity: 53.414 speed: 23490 wps

```

96 0.304 perplexity: 51.359 speed: 23501 wps
97 0.404 perplexity: 50.514 speed: 23483 wps
98 0.504 perplexity: 49.779 speed: 23483 wps
99 0.604 perplexity: 48.176 speed: 23496 wps
100 0.703 perplexity: 47.481 speed: 23490 wps
101 0.803 perplexity: 46.703 speed: 23473 wps
102 0.903 perplexity: 45.150 speed: 23460 wps
103 Epoch: 8 Train Perplexity: 44.260
104 Epoch: 8 Valid Perplexity: 120.259
105 Epoch: 9 Learning rate: 0.031
106 0.004 perplexity: 66.016 speed: 24117 wps
107 0.104 perplexity: 47.151 speed: 22937 wps
108 0.204 perplexity: 51.539 speed: 22904 wps
109 0.304 perplexity: 49.519 speed: 22993 wps
110 0.404 perplexity: 48.710 speed: 23037 wps
111 0.504 perplexity: 48.002 speed: 23073 wps
112 0.604 perplexity: 46.448 speed: 23068 wps
113 0.703 perplexity: 45.756 speed: 23065 wps
114 0.803 perplexity: 44.982 speed: 23085 wps
115 0.903 perplexity: 43.462 speed: 23086 wps
116 Epoch: 9 Train Perplexity: 42.586
117 Epoch: 9 Valid Perplexity: 120.535
118 Epoch: 10 Learning rate: 0.016
119 0.004 perplexity: 64.642 speed: 23044 wps
120 0.104 perplexity: 46.198 speed: 22610 wps
121 0.204 perplexity: 50.520 speed: 22409 wps
122 0.304 perplexity: 48.504 speed: 22626 wps
123 0.404 perplexity: 47.706 speed: 22884 wps
124 0.504 perplexity: 47.015 speed: 23025 wps
125 0.604 perplexity: 45.492 speed: 22952 wps
126 0.703 perplexity: 44.809 speed: 22868 wps
127 0.803 perplexity: 44.042 speed: 22834 wps
128 0.903 perplexity: 42.541 speed: 22901 wps
129 Epoch: 10 Train Perplexity: 41.671
130 Epoch: 10 Valid Perplexity: 120.385
131 Epoch: 11 Learning rate: 0.008
132 0.004 perplexity: 63.739 speed: 23199 wps
133 0.104 perplexity: 45.621 speed: 23363 wps
134 0.204 perplexity: 49.911 speed: 23356 wps
135 0.304 perplexity: 47.912 speed: 23224 wps
136 0.404 perplexity: 47.124 speed: 23025 wps
137 0.504 perplexity: 46.446 speed: 23045 wps
138 0.604 perplexity: 44.942 speed: 23097 wps
139 0.703 perplexity: 44.270 speed: 23078 wps
140 0.803 perplexity: 43.512 speed: 22958 wps
141 0.903 perplexity: 42.023 speed: 22952 wps
142 Epoch: 11 Train Perplexity: 41.159
143 Epoch: 11 Valid Perplexity: 120.067
144 Epoch: 12 Learning rate: 0.004
145 0.004 perplexity: 63.187 speed: 23541 wps
146 0.104 perplexity: 45.286 speed: 23184 wps
147 0.204 perplexity: 49.561 speed: 23173 wps
148 0.304 perplexity: 47.580 speed: 23181 wps
149 0.404 perplexity: 46.801 speed: 23073 wps

```

```

150 0.504 perplexity: 46.130 speed: 23139 wps
151 0.604 perplexity: 44.639 speed: 23209 wps
152 0.703 perplexity: 43.972 speed: 23263 wps
153 0.803 perplexity: 43.221 speed: 23306 wps
154 0.903 perplexity: 41.740 speed: 23336 wps
155 Epoch: 12 Train Perplexity: 40.881
156 Epoch: 12 Valid Perplexity: 119.783
157 Epoch: 13 Learning rate: 0.002
158 0.004 perplexity: 62.876 speed: 23633 wps
159 0.104 perplexity: 45.096 speed: 23555 wps
160 0.204 perplexity: 49.362 speed: 23525 wps
161 0.304 perplexity: 47.394 speed: 23554 wps
162 0.404 perplexity: 46.623 speed: 23563 wps
163 0.504 perplexity: 45.958 speed: 23584 wps
164 0.604 perplexity: 44.475 speed: 23578 wps
165 0.703 perplexity: 43.813 speed: 23566 wps
166 0.803 perplexity: 43.065 speed: 23570 wps
167 0.903 perplexity: 41.589 speed: 23577 wps
168 Epoch: 13 Train Perplexity: 40.733
169 Epoch: 13 Valid Perplexity: 119.608
170 Test Perplexity: 115.036
171 Test Loss: 4.745

```

The losses from the text files are plotted with the code below

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3 from matplotlib.ticker import MaxNLocator
4
5
6
7 def get_csv_data(filename):
8     return np.loadtxt(filename, dtype='float32', delimiter=',')
9
10
11 def plot_csv(filename, y_label, max_or_min):
12     data = get_csv_data(filename)
13     epochs = data[:,0]
14     acc = data[:,1]
15     string = ''
16     if (max_or_min == 'max'):
17         index = np.argmax(data[:,1])
18         string = 'Maximum'
19     else:
20         index = np.argmin(data[:,1])
21         string = 'Minimum'
22     # add plot and max/min point
23     plt.plot(epochs, acc, linestyle = '--', marker = 'o', color = 'b')
24     plt.plot(epochs[index], acc[index], marker = 'o', color='r')
25     # add vertical line
26     ax = plt.gca()
27     y_min, y_max = ax.get_ylim()
28     plt.vlines(epochs[index], y_min, acc[index], linestyle='dashed')
29     plt.ylim((y_min, y_max))

```

```

30 # set plot title
31 plt.title('{} vs Epoch'.format(y_label))
32 # delete x axis ticks that are within one of the max epoch
33 x_tick_arr = np.array(list(ax.get_xticks())).astype(int)
34 close_indices = np.where(np.logical_or(x_tick_arr==np.int(epochs[index] - 1),
    x_tick_arr==np.int(epochs[index] + 1)))[0]
35 ax.set_xticks(np.delete(x_tick_arr, close_indices))
36 # add max index to x axis
37 ax.set_xticks(list(ax.get_xticks()) + [epochs[index]])
38 # set plot x/y labels
39 plt.xlabel('Epoch')
40 plt.ylabel('{}'.format(y_label))
41 # add legend
42 string = string + ' ' + y_label + ': ' + str(acc[index])
43 plt.legend(['_nolegend_', string], loc = 'best')
44 # make the max point red
45 x_min, x_max = ax.get_xlim()
46 plt.xlim((0, epochs[-1:] + 1))
47 ax.get_xticklabels()[-1].set_color('red')
48
49 plt.show()
50
51
52 plot_csv('epoch_loss_train.txt', 'Training Loss', 'min')
53 plot_csv('epoch_loss_val.txt', 'Validation Loss', 'min')

```

Calling plotter.py gives the plots below






Figure 1: Training Loss vs Epoch

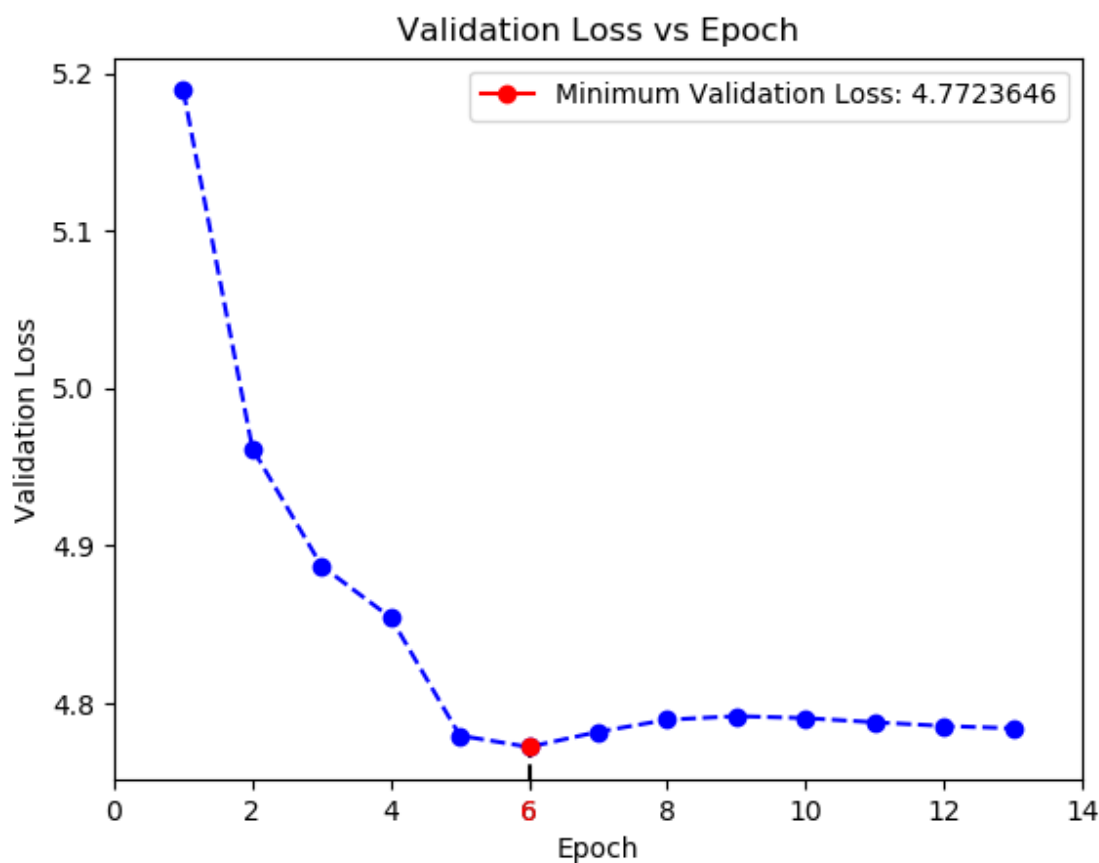


Figure 2: Validation Loss vs Epoch

The model is clearly overfitting, since the training perplexity and loss are far below that of the validation set. This could be solved by using a larger model, but was not used here due to time constraints.

### c. Test Results

It can be noted from the output of the training code gives a loss of 4.745, which is right around the validation loss, showing that the validation set is a good representation of the test set. In terms of perplexity, the test set achieved 115, the validation set 119, and the training set 40.733.

### d. Predicting Sentences with Four Different Words

The code for this section is below

```

1 # Copyright 2015 The TensorFlow Authors. All Rights Reserved.
2 #
3 # Licensed under the Apache License, Version 2.0 (the "License");
4 # you may not use this file except in compliance with the License.
5 # You may obtain a copy of the License at
6 #

```

```

7 # http://www.apache.org/licenses/LICENSE-2.0
8 #
9 # Unless required by applicable law or agreed to in writing, software
10 # distributed under the License is distributed on an "AS IS" BASIS,
11 # WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 # See the License for the specific language governing permissions and
13 # limitations under the License.
14 # =====
15
16 """Example / benchmark for building a PTB LSTM model.
17 Trains the model described in:
18 (Zaremba, et. al.) Recurrent Neural Network Regularization
19 http://arxiv.org/abs/1409.2329
20 There are 3 supported model configurations:
21 =====
22 | config | epochs | train | valid | test
23 =====
24 | small | 13 | 37.99 | 121.39 | 115.91
25 | medium | 39 | 48.45 | 86.16 | 82.07
26 | large | 55 | 37.87 | 82.62 | 78.29
27 The exact results may vary depending on the random initialization.
28 The hyperparameters used in the model:
29 - init_scale - the initial scale of the weights
30 - learning_rate - the initial value of the learning rate
31 - max_grad_norm - the maximum permissible norm of the gradient
32 - num_layers - the number of LSTM layers
33 - num_steps - the number of unrolled steps of LSTM
34 - hidden_size - the number of LSTM units
35 - max_epoch - the number of epochs trained with the initial learning rate
36 - max_max_epoch - the total number of epochs for training
37 - keep_prob - the probability of keeping weights in the dropout layer
38 - lr_decay - the decay of the learning rate for each epoch after "max_epoch"
39 - batch_size - the batch size
40 - rnn_mode - the low level implementation of lstm cell: one of CUDNN,
41 BASIC, or BLOCK, representing cudnn_lstm, basic_lstm, and
42 lstm_block_cell classes.
43 The data required for this example is in the data/ dir of the
44 PTB dataset from Tomas Mikolov's webpage:
45 $ wget http://www.fit.vutbr.cz/~imikolov/rnnlm/simple-examples.tgz
46 $ tar xvf simple-examples.tgz
47 To run:
48 $ python ptb_word_lm.py --data_path=simple-examples/data/
49 """
50 from __future__ import absolute_import
51 from __future__ import division
52 from __future__ import print_function
53
54 import time
55
56 import numpy as np
57 import tensorflow as tf
58
59 import reader
60 import util

```

```

61
62 from tensorflow.python.client import device_lib
63
64 flags = tf.flags
65 logging = tf.logging
66
67 flags.DEFINE_string(
68     "model", "small",
69     "A type of model. Possible options are: small, medium, large.")
70 flags.DEFINE_string("data_path", None,
71     "Where the training/test data is stored.")
72 flags.DEFINE_string("save_path", None,
73     "Model output directory.")
74 flags.DEFINE_bool("use_fp16", False,
75     "Train using 16-bit floats instead of 32bit floats")
76 flags.DEFINE_integer("num_gpus", 1,
77     "If larger than 1, Grappler AutoParallel optimizer "
78     "will create multiple training replicas with each GPU "
79     "running one replica.")
80 flags.DEFINE_string("rnn_mode", None,
81     "The low level implementation of lstm cell: one of CUDNN, "
82     "BASIC, and BLOCK, representing cudnn_lstm, basic_lstm, "
83     "and lstm_block_cell classes.")
84 FLAGS = flags.FLAGS
85 BASIC = "basic"
86 CUDNN = "cudnn"
87 BLOCK = "block"
88
89
90 def data_type():
91     return tf.float16 if FLAGS.use_fp16 else tf.float32
92
93
94 class PTBInput(object):
95     """The input data."""
96
97     def __init__(self, config, data, name=None):
98         self.batch_size = batch_size = config.batch_size
99         self.num_steps = num_steps = config.num_steps
100         self.epoch_size = ((len(data) // batch_size) - 1) // num_steps
101         self.input_data, self.targets = reader.ptb_producer(
102             data, batch_size, num_steps, name=name)
103
104
105 class PTBModel(object):
106     """The PTB model."""
107
108     def __init__(self, is_training, config, input_):
109         self._is_training = is_training
110         # self._input = input_
111         self._rnn_params = None
112         self._cell = None
113         self.batch_size = batch_size = config.batch_size
114         self.num_steps = num_steps = config.num_steps

```

```

115 size = config.hidden_size
116 vocab_size = config.vocab_size
117
118 self._input_data = tf.placeholder(tf.int32, [batch_size, num_steps])
119 self._targets = tf.placeholder(tf.int32, [batch_size, num_steps])
120
121 with tf.device("/cpu:0"):
122     embedding = tf.get_variable(
123         "embedding", [vocab_size, size], dtype=data_type())
124     inputs = tf.nn.embedding_lookup(embedding, self._input_data)
125
126 if is_training and config.keep_prob < 1:
127     inputs = tf.nn.dropout(inputs, config.keep_prob)
128
129 output, state = self._build_rnn_graph(inputs, config, is_training)
130
131 softmax_w = tf.get_variable(
132     "softmax_w", [size, vocab_size], dtype=data_type())
133 softmax_b = tf.get_variable("softmax_b", [vocab_size], dtype=data_type())
134 logits = tf.nn.xw_plus_b(output, softmax_w, softmax_b)
135 self.sample = tf.multinomial(logits, 1)
136 # Reshape logits to be a 3-D tensor for sequence loss
137 logits = tf.reshape(logits, [self.batch_size, self.num_steps, vocab_size])
138
139 # Use the contrib sequence loss and average over the batches
140 loss = tf.contrib.seq2seq.sequence_loss(
141     logits,
142     self._targets,
143     tf.ones([self.batch_size, self.num_steps], dtype=data_type()),
144     average_across_timesteps=False,
145     average_across_batch=True)
146
147 # Update the cost
148 self._cost = tf.reduce_sum(loss)
149 self._final_state = state
150
151 if not is_training:
152     return
153
154 self._lr = tf.Variable(0.0, trainable=False)
155 tvars = tf.trainable_variables()
156 grads, _ = tf.clip_by_global_norm(tf.gradients(self._cost, tvars),
157     config.max_grad_norm)
158 optimizer = tf.train.GradientDescentOptimizer(self._lr)
159 self._train_op = optimizer.apply_gradients(
160     zip(grads, tvars),
161     global_step=tf.train.get_or_create_global_step())
162
163 self._new_lr = tf.placeholder(
164     tf.float32, shape=[], name="new_learning_rate")
165 self._lr_update = tf.assign(self._lr, self._new_lr)
166
167 def _build_rnn_graph(self, inputs, config, is_training):
168     if config.rnn_mode == CUDNN:

```

```

169         return self._build_rnn_graph_cudnn(inputs, config, is_training)
170     else:
171         return self._build_rnn_graph_lstm(inputs, config, is_training)
172
173 def _build_rnn_graph_cudnn(self, inputs, config, is_training):
174     """Build the inference graph using CUDNN cell."""
175     inputs = tf.transpose(inputs, [1, 0, 2])
176     self._cell = tf.contrib.cudnn_rnn.CudnnLSTM(
177         num_layers=config.num_layers,
178         num_units=config.hidden_size,
179         input_size=config.hidden_size,
180         dropout=1 - config.keep_prob if is_training else 0)
181     params_size_t = self._cell.params_size()
182     self._rnn_params = tf.get_variable(
183         "lstm_params",
184         initializer=tf.random_uniform(
185             [params_size_t], -config.init_scale, config.init_scale),
186         validate_shape=False)
187     c = tf.zeros([config.num_layers, self.batch_size, config.hidden_size],
188                 tf.float32)
189     h = tf.zeros([config.num_layers, self.batch_size, config.hidden_size],
190                 tf.float32)
191     self._initial_state = (tf.contrib.rnn.LSTMStateTuple(h=h, c=c),)
192     outputs, h, c = self._cell(inputs, h, c, self._rnn_params, is_training)
193     outputs = tf.transpose(outputs, [1, 0, 2])
194     outputs = tf.reshape(outputs, [-1, config.hidden_size])
195     return outputs, (tf.contrib.rnn.LSTMStateTuple(h=h, c=c),)
196
197 def _get_lstm_cell(self, config, is_training):
198     if config.rnn_mode == BASIC:
199         return tf.contrib.rnn.BasicLSTMCell(
200             config.hidden_size, forget_bias=0.0, state_is_tuple=True,
201             reuse=not is_training)
202     if config.rnn_mode == BLOCK:
203         return tf.contrib.rnn.LSTMBlockCell(
204             config.hidden_size, forget_bias=0.0)
205     raise ValueError("rnn_mode %s not supported" % config.rnn_mode)
206
207 def _build_rnn_graph_lstm(self, inputs, config, is_training):
208     """Build the inference graph using canonical LSTM cells."""
209     # Slightly better results can be obtained with forget gate biases
210     # initialized to 1 but the hyperparameters of the model would need to be
211     # different than reported in the paper.
212     def make_cell():
213         cell = self._get_lstm_cell(config, is_training)
214         if is_training and config.keep_prob < 1:
215             cell = tf.contrib.rnn.DropoutWrapper(
216                 cell, output_keep_prob=config.keep_prob)
217         return cell
218
219     cell = tf.contrib.rnn.MultiRNNCell(
220         [make_cell() for _ in range(config.num_layers)], state_is_tuple=True)
221
222     self._initial_state = cell.zero_state(config.batch_size, data_type())

```

```

223 state = self._initial_state
224 # Simplified version of tf.nn.static_rnn().
225 # This builds an unrolled LSTM for tutorial purposes only.
226 # In general, use tf.nn.static_rnn() or tf.nn.static_state_saving_rnn().
227 #
228 # The alternative version of the code below is:
229 #
230 inputs = tf.unstack(inputs, num=self.num_steps, axis=1)
231 outputs, state = tf.nn.static_rnn(cell, inputs,
232                                   initial_state=self._initial_state)
233 # outputs = []
234 # with tf.variable_scope("RNN"):
235 #     for time_step in range(self.num_steps):
236 #         if time_step > 0: tf.get_variable_scope().reuse_variables()
237 #         (cell_output, state) = cell(inputs[:, time_step, :], state)
238 #         outputs.append(cell_output)
239 output = tf.reshape(tf.concat(outputs, 1), [-1, config.hidden_size])
240 return output, state
241
242 def assign_lr(self, session, lr_value):
243     session.run(self._lr_update, feed_dict={self._new_lr: lr_value})
244
245 def export_ops(self, name):
246     """Exports ops to collections."""
247     self._name = name
248     ops = {util.with_prefix(self._name, "cost"): self._cost}
249     if self._is_training:
250         ops.update(lr=self._lr, new_lr=self._new_lr, lr_update=self._lr_update)
251         if self._rnn_params:
252             ops.update(rnn_params=self._rnn_params)
253     for name, op in ops.items():
254         tf.add_to_collection(name, op)
255     self._initial_state_name = util.with_prefix(self._name, "initial")
256     self._final_state_name = util.with_prefix(self._name, "final")
257     util.export_state_tuples(self._initial_state, self._initial_state_name)
258     util.export_state_tuples(self._final_state, self._final_state_name)
259
260 def import_ops(self):
261     """Imports ops from collections."""
262     if self._is_training:
263         self._train_op = tf.get_collection_ref("train_op")[0]
264         self._lr = tf.get_collection_ref("lr")[0]
265         self._new_lr = tf.get_collection_ref("new_lr")[0]
266         self._lr_update = tf.get_collection_ref("lr_update")[0]
267         rnn_params = tf.get_collection_ref("rnn_params")
268         if self._cell and rnn_params:
269             params_saveable = tf.contrib.cudnn_rnn.RNNParamsSaveable(
270                 self._cell,
271                 self._cell.params_to_canonical,
272                 self._cell.canonical_to_params,
273                 rnn_params,
274                 base_variable_scope="Model/RNN")
275             tf.add_to_collection(tf.GraphKeys.SAVEABLE_OBJECTS, params_saveable)
276         self._cost = tf.get_collection_ref(util.with_prefix(self._name, "cost"))[0]

```

```

277     num_replicas = FLAGS.num_gpus if self._name == "Train" else 1
278     self._initial_state = util.import_state_tuples(
279         self._initial_state, self._initial_state_name, num_replicas)
280     self._final_state = util.import_state_tuples(
281         self._final_state, self._final_state_name, num_replicas)
282
283
284     @property
285     def initial_state(self):
286         return self._initial_state
287
288     @property
289     def cost(self):
290         return self._cost
291
292     @property
293     def final_state(self):
294         return self._final_state
295
296     @property
297     def lr(self):
298         return self._lr
299
300     @property
301     def train_op(self):
302         return self._train_op
303     @property
304     def input_data(self):
305         return self._input_data
306
307     @property
308     def targets(self):
309         return self._targets
310     @property
311     def initial_state_name(self):
312         return self._initial_state_name
313
314     @property
315     def final_state_name(self):
316         return self._final_state_name
317
318
319
320 class SmallConfig(object):
321     """Small config."""
322     init_scale = 0.1
323     learning_rate = 1.0
324     max_grad_norm = 5
325     num_layers = 2
326     num_steps = 20
327     hidden_size = 200
328     max_epoch = 4
329     max_max_epoch = 13
330     keep_prob = 1.0

```



```

331 lr_decay = 0.5
332 batch_size = 20
333 vocab_size = 10000
334 rnn_mode = BLOCK
335
336
337 class MediumConfig(object):
338     """Medium config."""
339     init_scale = 0.05
340     learning_rate = 1.0
341     max_grad_norm = 5
342     num_layers = 2
343     num_steps = 35
344     hidden_size = 650
345     max_epoch = 6
346     max_max_epoch = 39
347     keep_prob = 0.5
348     lr_decay = 0.8
349     batch_size = 20
350     vocab_size = 10000
351     rnn_mode = BLOCK
352
353
354 class LargeConfig(object):
355     """Large config."""
356     init_scale = 0.04
357     learning_rate = 1.0
358     max_grad_norm = 10
359     num_layers = 2
360     num_steps = 35
361     hidden_size = 1500
362     max_epoch = 14
363     max_max_epoch = 55
364     keep_prob = 0.35
365     lr_decay = 1 / 1.15
366     batch_size = 20
367     vocab_size = 10000
368     rnn_mode = BLOCK
369
370
371 class TestConfig(object):
372     """Tiny config, for testing."""
373     init_scale = 0.1
374     learning_rate = 1.0
375     max_grad_norm = 1
376     num_layers = 1
377     num_steps = 2
378     hidden_size = 2
379     max_epoch = 1
380     max_max_epoch = 1
381     keep_prob = 1.0
382     lr_decay = 0.5
383     batch_size = 20
384     vocab_size = 10000

```

```

385     rnn_mode = BLOCK
386
387
388 def run_epoch(session, model, eval_op=None, verbose=False):
389     """Runs the model on the given data."""
390     start_time = time.time()
391     costs = 0.0
392     iters = 0
393     state = session.run(model.initial_state)
394
395     fetches = {
396         "cost": model.cost,
397         "final_state": model.final_state,
398     }
399     if eval_op is not None:
400         fetches["eval_op"] = eval_op
401
402     for step in range(model.input.epoch_size):
403         feed_dict = {}
404         for i, (c, h) in enumerate(model.initial_state):
405             feed_dict[c] = state[i].c
406             feed_dict[h] = state[i].h
407
408         vals = session.run(fetches, feed_dict)
409         cost = vals["cost"]
410         state = vals["final_state"]
411
412         costs += cost
413         iters += model.input.num_steps
414
415         if verbose and step % (model.input.epoch_size // 10) == 10:
416             print("%.3f perplexity: %.3f speed: %.0f wps" %
417                   (step * 1.0 / model.input.epoch_size, np.exp(costs / iters),
418                    iters * model.input.batch_size * max(1, FLAGS.num_gpus) /
419                    (time.time() - start_time)))
420
421     return np.exp(costs / iters), (costs / iters), logits
422
423
424 def get_config():
425     """Get model config."""
426     config = None
427     if FLAGS.model == "small":
428         config = SmallConfig()
429     elif FLAGS.model == "medium":
430         config = MediumConfig()
431     elif FLAGS.model == "large":
432         config = LargeConfig()
433     elif FLAGS.model == "test":
434         config = TestConfig()
435     else:
436         raise ValueError("Invalid model: %s", FLAGS.model)
437     if FLAGS.rnn_mode:
438         config.rnn_mode = FLAGS.rnn_mode

```

```

439     if FLAGS.num_gpus != 1 or tf.__version__ < "1.3.0" :
440         config.rnn_mode = BASIC
441     return config
442
443 def get_sentence(session, model, data, num_samples):
444     # get initial state
445     samples = []
446     state = session.run(model.initial_state)
447     # fetches are the final state and the prediction
448     fetches = [model.final_state, model.sample]
449     sample = None
450     # get the next word, can be from multiple words in data
451     for x in data:
452         feed_dict = {}
453         feed_dict[model.input_data] = [[x]]
454         for layer_num, (c, h) in enumerate(model.initial_state):
455             feed_dict[c] = state[layer_num].c
456             feed_dict[h] = state[layer_num].h
457
458         state, sample = session.run(fetches, feed_dict)
459
460     # append next word
461     samples.append(sample[0][0])
462
463     # for num_samples, append words
464     k = 1
465     while k < num_samples:
466         feed_dict = {}
467         feed_dict[model.input_data] = [[samples[-1]]]
468         for layer_num, (c, h) in enumerate(model.initial_state):
469             feed_dict[c] = state[layer_num].c
470             feed_dict[h] = state[layer_num].h
471         state, sample = session.run(fetches, feed_dict)
472         samples.append(sample[0][0])
473         # if sample is <eos>, return samples early
474         if (sample[0][0] == 2):
475             return samples
476
477         k += 1
478     return samples
479
480 def print_sentence(items, id_to_word):
481     sentence = ''
482     for item in items:
483         sentence = sentence + ' ' + id_to_word[item]
484     return sentence
485
486 def main(_):
487     if not FLAGS.data_path:
488         raise ValueError("Must set --data_path to PTB data directory")
489     gpus = [
490         x.name for x in device_lib.list_local_devices() if x.device_type == "GPU"
491     ]

```

```

493 if FLAGS.num_gpus > len(gpus):
494     raise ValueError(
495         "Your machine has only %d gpus "
496         "which is less than the requested --num_gpus=%d."
497         % (len(gpus), FLAGS.num_gpus))
498
499 raw_data = reader.ptb_raw_data(FLAGS.data_path)
500 train_data, valid_data, test_data, _, word_to_id = raw_data
501 id_to_word = dict(zip(word_to_id.values(), word_to_id.keys()))
502
503 config = get_config()
504 eval_config = get_config()
505 eval_config.batch_size = 1
506 eval_config.num_steps = 1
507
508 with tf.Graph().as_default():
509     initializer = tf.random_uniform_initializer(-config.init_scale,
510                                                config.init_scale)
511
512     with tf.name_scope("Train"):
513         train_input = PTBInput(config=config, data=train_data, name="TrainInput")
514         with tf.variable_scope("Model", reuse=None, initializer=initializer):
515             m = PTBModel(is_training=True, config=config, input_=train_input)
516             tf.summary.scalar("Training Loss", m.cost)
517             tf.summary.scalar("Learning Rate", m.lr)
518
519
520     with tf.name_scope("Test"):
521         test_input = PTBInput(
522             config=eval_config, data=test_data, name="TestInput")
523         with tf.variable_scope("Model", reuse=True, initializer=initializer):
524             mtest = PTBModel(is_training=False, config=eval_config,
525                             input_=test_input)
526
527     saver = tf.train.Saver()
528     titan_v = "0"
529     titan_x = "1"
530     gpu_options = tf.GPUOptions(visible_device_list=titan_v)
531     config = tf.ConfigProto(gpu_options=gpu_options)
532     sv = tf.train.Supervisor(logdir=FLAGS.save_path)
533     with sv.managed_session() as session:
534         sv.saver.restore(session, './ptb_lstm.ckpt')
535         for i in range(1):
536             user_input = input("Enter your starting words: ")
537             max_num_words = int(input("Max words: "))
538             starting_words = user_input.split()
539             print("Starting Words: %s" % print_sentence([word_to_id[x] for x in
540                                                         starting_words], id_to_word))
541             sentence = get_sentence(session, mtest, [word_to_id[word] for word in
542                                                         starting_words], max_num_words)
543             print("Resulting Sentence: %s" % print_sentence(sentence, id_to_word))
544

```

```

545
546 if __name__ == "__main__":
547     tf.app.run()

```

Note the main changes are using placeholders as input instead of hard coded lengths, and the `get_sentence` and `print_sentence` methods. The code goes for 20 words or the end of a sentence. The results for four chosen words are below

### Word 1: the

```

1 Enter your starting words: the
2 Max words: 20
3 Starting Words: the
4 Resulting Sentence: <unk> of america and the <unk> few face other than the absolutely
   <unk> of the outcome of a new post

```

### Word 2: words

```

1 Enter your starting words: words
2 Max words: 20
3 Starting Words: words
4 Resulting Sentence: are clearly considered more often than saying it would be worth
   <unk> <eos>

```

### Word 3: software

```

1 Enter your starting words: software
2 Max words: 20
3 Starting Words: software
4 Resulting Sentence: factory policies sweat lower limited and pricing <eos>

```

### Word 4: murder

```

1 Enter your starting words: murder
2 Max words: 20
3 Starting Words: murder
4 Resulting Sentence: a private championship to buy them the championship <unk> available
   in a major market <eos>

```

## e. Predicting Sentences with given words

### Word 1: north

```

1 Enter your starting words: north
2 Max words: 20
3 Starting Words: north

```

4 Resulting Sentence: american pacific corp. and <unk> in california contributed to this  
london options and mark <eos>

### Word 2: wall

1 Enter your starting words: wall  
2 Max words: 20  
3 Starting Words: wall  
4 Resulting Sentence: street 's biggest investment investment firms led the identity field  
to farmers in a televised <unk> <unk> <eos>

### Word 3: truth

1 Enter your starting words: truth  
2 Max words: 20  
3 Starting Words: truth  
4 Resulting Sentence: dr. freeman says when an <unk> imposed a significant N portion of  
the award <eos>

### Word 4: california

1 Enter your starting words: california  
2 Max words: 20  
3 Starting Words: california  
4 Resulting Sentence: may not give the profile of a massive dollar without their own foot  
<eos>

### Word 5: health

1 Enter your starting words: health  
2 Max words: 20  
3 Starting Words: health  
4 Resulting Sentence: insurance operation <eos>